

Evaluating the Performance of a Stacked Long Short-Term Memory Deep Learning model on AEX Stock Price Prediction

Jeff Maes, Daan Cuppen, Isaac Seminck, Benjamin CD De Bosscher, Kipras Paliuis
Delft University of Technology

Faculty of Computer Science

Group 55

<https://github.com/benjamindeboscher/DL-LSTM-Stock-Prediction>

Abstract

Recent trends have shown the successful application of deep learning architectures on stock data. In most cases, recurrent neural network models are used for this form of time series data. In this midterm paper, an existing long short-term memory (LSTM) model will be further developed by adding the volume as an extra input dimension. The improved model will be analysed by means of several key performance indicators and graphs. From the preliminary visual results, this proves to be improving the prediction accuracy of the model.

LSTM model only considers previous stock data and thus no other influences on the stock price such as news events, quarterly company results, macro economic variables, etc. will be considered. In this way, we hope to achieve better results which can be used for longer periods of time as we do not want to retrain our model every trading day. The next steps of the research for the final article are described in the reflection (section 5). The research question can be stated as follows: How can the prediction performance of an LSTM model on the stock prices of companies in the Amsterdam Exchange Index (AEX) be improved?

1. Introduction

In stock prediction, the goal is predicting future trends of a company stock as doing this successfully could yield a significant profit or a reduction in losses. An algorithm capable of doing this could help a trader during the decision-making process of buying, keeping or selling stocks. Stock prices have been predicted in the past using techniques such as moving averages, however, this technique has several drawbacks, as will be elaborated upon in section 3.

In this paper, modifications are made to an existing stock price prediction model based on a long short-term memory deep learning model [1] to see if its performance can be increased. Currently, this model only considers the effect of stock price changes, but research suggests that also the volume of a stock can be of interest to include in the deep model [5]. In short, the volume of a stock is how many stocks have been traded on a single trading day. A significant change in volume can reflect stock volatility, which refers to a drastic decrease or increase in stock value, causing an imbalance in trade orders. The second alteration that is made to the algorithm is to vary the prediction sequence lengths. In this way, the correlation between prediction performance and prediction length sequence. Furthermore, the

2. Problem Statement

Accurately predicting stock data is a difficult task. This paper aims to improve performance of an already existing deep algorithm. The data set on which the algorithm is applied, are the stock data of AEX companies. Obtaining the relevant stock data is not a problem as it is widely available. It can be downloaded but it can also be imported using an API (eg. Quandl), etc. For this project, downloading the data is preferred as this makes it possible to go back further in time which results in more data. The data used for this project are all the underlying stocks of the Amsterdam Exchange Index. For every trading day, the open, close, low and high are obtained, so is the traded volume.

The approach taken in this paper is to train an LSTM-model on all stocks in the AEX index, which takes into account both mid-price and volume. The performance is checked after every alteration performed on the algorithm. Therefore, the effect of each modification can be analysed and the optimal configuration for the model can be determined. To specify in more detail which 25 stocks are analysed, a summary is given in Table 1.

AALB	ASML	GTO	MT	REN
ABN	ASRN	HEIA	NN	UNA
AD	ATC	INGA	PHIA	URW
AGN	DSM	KPN	RAND	VPK
AKZA	GPLG	LIGHT	RDSA	WKL

Table 1. Ticker symbols of the investigated AEX stocks.

3. Method

To predict a future trend of a stock, there are several methods. One could for example opt for a well-known moving average technique. But there are also other techniques available such as a momentum-based algorithm, an LSTM model, etc. In this section, there will be elaborated on the method that is used for the stock prediction model.

3.1. Moving Average

Moving average techniques are widely applied for several prediction purposes; they give good results [1]. However, one should look further than just the optical illusion they actually might be. These techniques come with subtle but significant risks to investors [4].

First of all, they cannot predict more than one step in the future. The same prediction comes back for further predictions. This is a huge drawback for investors; instead of knowing the stock price for the next step, they would like to know the future trend (whether the stock will rise or fall) for more than one prediction step. Furthermore as their name says, moving average techniques use an average of their past values. This can be simple, exponential, weighted, cumulative, etc. Because of this average, they have difficulties to cope with high volatile stocks. Then there is also the time period. Choosing a certain time period influences the general trend a stock typifies. This means that a moving average prediction might include a certain undesired bias. Lastly, stocks show often an oscillatory pattern. This effect is hidden when using moving average methods.

All this together makes moving average techniques ideal to analyse the past behaviour of a stock. However in order to do future predictions, it lacks some effectiveness. There are better methods, such as a long short-term memory (LSTM) deep learning model.

3.2. LSTM Model

An LSTM model is very powerful method to model sequence dependent input. They are classified under the bigger group of Recurrent Neural Networks, RNNs in short. RNNs are able to persist information as they consist of networks with loops within them [2]. The general principle of an LSTM is visualised on the left side of Figure 1. The x stands for the input of the LSTM-cell A which outputs h . The right side of Figure 1 displays the structure of the stacked LSTM model that is taken as the basis model in this

paper. The basis model has three LSTM-cells per time step and is able to predict multiple time steps in the future. In Figure 1, the number of hidden nodes between the LSTM-cells per time step are indicated and equal 200 for the two hidden layers, and 150 for the regression layer.

As can be seen from the visualisation; the loops pass information throughout the steps. This means that this architecture has the ability to remember information. This remembering property is exactly the reason why LSTMs (RNNs in general) are much more powerful than other deep learning architectures (e.g. a convolutional neural network, CNN in short) to model time series data.




Figure 1. Visualisation of a stacked LSTM [3].

Next to remembering, LSTM models also have the ability to forget information with the help of a sigmoid layer. This property made them eliminate the long-term dependencies problem conventional RNNs traditionally struggle with [3]. The problem with these long-term dependencies is that the gradients might either vanish or explode; a consequence which should be avoided. The general outline of an LSTM network is visualised in Figure 2.

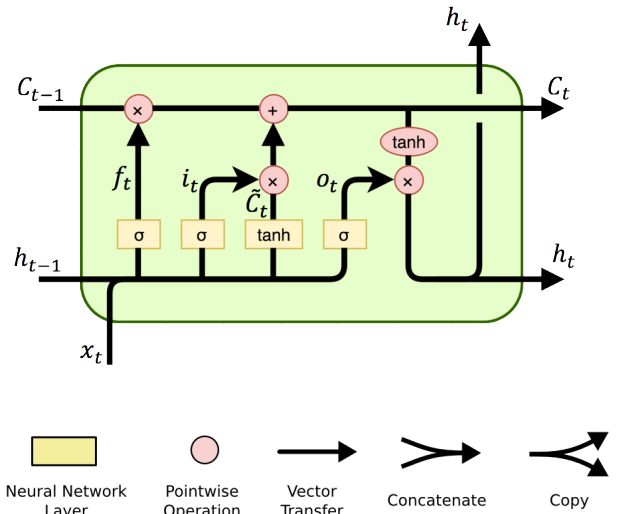


Figure 2. Visualisation of an LSTM cell [3].

Readily, five fundamental components can be identified [1]. First, there is the cell state c_t which represents the memory of the cell. Secondly, an LSTM consist of a hidden state h_t . Essentially, the hidden state contains the output which is merely a filtered version of the cell state. Furthermore, the input gate i_t can be identified. It decides how much of the input x_t flows to the cell state c_t . Fourthly, there is a forget gate f_t which determines the amount of information that flows from the previous hidden state h_{t-1} and the input x_t to the current cell state c_t . Lastly, the output gate o_t decides the amount of information that flows from the cell state c_t to the hidden state h_t .

To summarise, an LSTM model is a powerful tool in modelling time series, The basis LSTM model in this paper has three LSTM-cells per time step, is able to predict multiple steps into the future and is a SISO-system.

3.3. Data Preprocessing

During this midterm paper, only one stock (PHIA) will be considered. For this stock, the total number of data inputs is slightly over 6000 trading days. This input data was first split into a training set of 5000 data points and a test set of over 1000 data points. Then both the test and training data sets were normalised with respect to the training set, because the algorithm should have no access to the test data in advance. Due to the fact that a particular stock increases substantially over the years, different time periods have different value ranges. Therefore the data is normalised by splitting the full series of data into four normalising windows. This introduces a break at the end of every window, thus four data points will be lost due to this approach of normalisation. For the PHIA stock this normalising window was chosen to be 1000 datapoints. Subsequently, the training data was smoothed out by using the exponential moving average in order to erase the sharpness of the prices. This helps the LSTM to train effectively and obtain improved performance as the LSTM is trained only to predict the general trend of the stock price.

3.4. Data Augmentation

Normally the prediction output \hat{y}_i is based upon one input x_i . In order to augment the data, the input is sampled from the full input sequence $[x_{t+1}, x_{t+2}, \dots, x_{t+N}]$ where N is the number of data points, resulting in input and output batches. Here, the assumption is made that the price between two subsequent input data points does not change much. This is a reasonable assumption for stock prices. In Figure 3, this data augmentation is explained visually. In a batch each entry is taken to be the first data point after a cursor. Thus the first sampled input batch is $[x_0, x_{N/b}, \dots, x_{N-b}]$. The cursors start at the beginning of a new segment in the input sequence, which is N/b data points long. For the second batch the cursors

are shifted to the right by one array element, resulting in $[x_1, x_{N/b+1}, \dots, x_{N-b+1}]$. This means that the data augmentation results in b batches, each N/b long. The batch size is chosen to be 500 with a input training data sequence of 5000.




Figure 3. Visualisation of data augmentation [1].

3.5. Adding Volume

The original LSTM model [1] is a stacked SISO (single input single output) LSTM model with three cells per time step and is able to predict multiple time steps in the future.

The original model only takes the mid-price as input and output. The volume is added as an additional input to this model to see if stock price prediction performance increases and the input thus becomes two-dimensional. In the prediction phase, the input data for the model is the output of the previous time step. Therefore, the model also needs to output volume when predicting multiple time steps into the future. If the model would only predict one time step into the future, the volume output would be unnecessary.

3.6. Varying Prediction Length

Another way of improving the current LSTM model is by looking into the effect of varying the length of time for which predictions are made. Ideally, the model would predict for a longer time window, making it usable for longer periods of time without the need for retraining and allowing for more long-term predictions. However, using a longer prediction window is expected to have negative effects on their accuracy. It is thus worth looking into the effects on the performance by varying the prediction window length and possibly determining the optimal window length allowing for the ideal prediction length to model performance ratio.

3.7. Performance Indicators

In order to evaluate the performance of the LSTM more accurately, the following performance indicators are used. The mean absolute error (MAE) measures the average magnitude of the errors in a set of predictions, without taking

their direction into consideration. The MAE is defined as follows:

$$\text{MAE} = \frac{1}{n} \sum_{j=1}^n |y_j - \hat{y}_j| \quad (1)$$

Where n is the number of prediction steps (number of unrollings of the LSTM), \hat{y}_j are the predictions and y_j are the mid-prices. The root mean squared error (RMSE) is defined as:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{j=1}^n (y_j - \hat{y}_j)^2} \quad (2)$$

The RMSE gives a higher weight to large errors compared to the MAE (same weight to all errors) as the errors are squared before averaged. The linear correlation coefficient (LCC) measures the strength of the linear relationship between two variables. The LCC for a sample is defined as:

$$r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}} \quad (3)$$

Where \bar{x} and \bar{y} are the means of the predictions and the mid-prices over the number of prediction step. Essentially, this equation is equal to: $\rho_{X,Y} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y}$. Lastly the maximum absolute error (MAE) measures the maximum error and indicates how the predictions follow extreme trends in the stock price.

4. Results

We already achieved adding the stock volume next to the stock price as an input variables for the network to train on. The results of this extra inputs are compared with tables 3 and 2. It can be easily seen that the performance indicators do not change significantly after the fifth epoch. Comparing tables 3 and 2, the LSTM excluding volume performs marginally better across all performance indicators.

	MAE	MRE	RMSE	LC
epoch 1	0.5849	0.5728	0.4234	0.5476
epoch 5	0.1098	0.1135	0.0854	2.4416
epoch 10	0.1140	0.1199	0.08912	2.2471
epoch 15	0.1125	0.1185	0.0879	4.887
epoch 20	0.1107	0.1174	0.0863	2.8042
epoch 25	0.1101	0.1171	0.0858	2.3500
epoch 30	0.1113	0.1178	0.0868	3.2977

Table 2. Results for different performance indicators for LSTM excluding volume.

	MAE	MRE	RMSE	LC
epoch 1	0.8304	0.7937	0.5952	0.5857
epoch 5	0.1058	0.1096	0.0824	1.2686
epoch 10	0.1415	0.1319	0.1102	0.6968
epoch 15	0.1230	0.1185	0.0959	1.3257
epoch 20	0.1221	0.1180	0.0952	1.2834
epoch 25	0.1228	0.1186	0.0958	1.2172
epoch 30	0.1234	0.1189	0.0963	1.1571

Table 3. Results for different performance indicators for LSTM including volume.

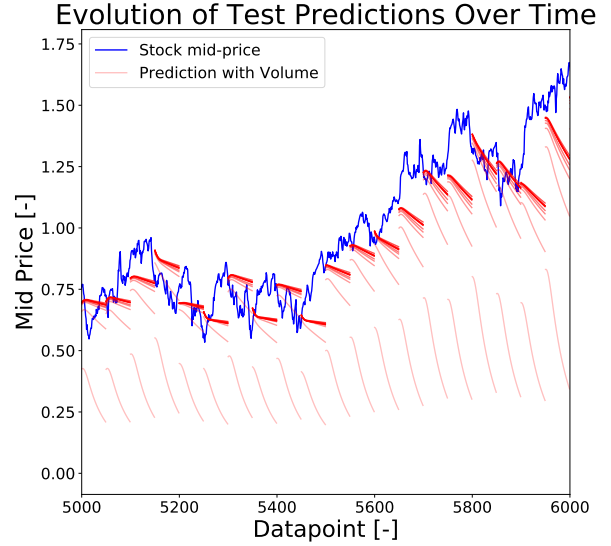


Figure 4. Evolution of the mid-price predictions over time for the model with volume, for the Philips (PHIA) stock.




Figure 5. Prediction of mid-price for the model with and without volume for the Philips (PHIA) stock.

In Figure 4 the direct outputs (all output sequences) are shown of the LSTM with volume. From this figure it can be seen that the predictions always tend to go down. On top of this one prediction of the sequence initiates at the wrong place (lines at the bottom of the figure). We aim to fix this wrong initialisation by the final report.

In Figure 5, the prediction of the normalised mid-price is plotted for the 1,000 data points considered in the testing phase. Predictions are plotted for both models, thus with and without taking volume into account. It can be seen that the predictions with volume perform better than the model without volume, however, the results from Table 3 indicate that the opposite is true.

5. Reflection

Currently, we are right on track as was originally planned. It will thus be possible to finish the initial proposal. In the next stage of this project we aim to achieve the following:

1. **Saving the trained LSTM model.** This is convenient as the model does not have to be run to get the parameters of the LSTM model. The model can just be loaded from an already saved file.
2. **Training of the LSTM on all stocks.** Now the model is only trained on one stock. This is not much data. The goal is to train the model on all 25 stocks of the AEX in order to improve our predictions.
3. **Run the model multiple times to check error values.** The model is always initialised randomly, this affects the performance of the model. Therefore we want to run the same model several times and then average these results.
4. **Hyperparameter tuning.** In the current model the hyperparameters are not optimised. We want to perform hyperparameter tuning by using a random grid search.
5. If there is enough time we want to compare **different LSTM model structures** such as GRU (gated recurrent unit) and LSTM with peepholes.

References

- [1] Thushan Ganegedara. Stock market predictions with lstm in python. <https://www.datacamp.com/community/tutorials/lstm-python-stock-market>, May 3rd, 2018. Accessed on: 2019-20-05.
- [2] Andrej Karpathy. The unreasonable effectiveness of recurrent neural networks. <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>, May 21, 2015. Accessed on: 2019-26-05.
- [3] Christopher Olah. Understanding lstm networks. <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>, August 27, 2015. Accessed on: 2019-26-05.
- [4] Kalen Smith. The 7 pitfalls of moving averages. <https://www.investopedia.com/articles/trading/11/pitfalls-moving-averages.asp>, June 25th, 2018. Accessed on: 2019-25-05.
- [5] Tianyi Wang and Zhuo Huang. *The Relationship between Volatility and Trading Volume in the Chinese Stock Market: A Volatility Decomposition Perspective*. 2012.